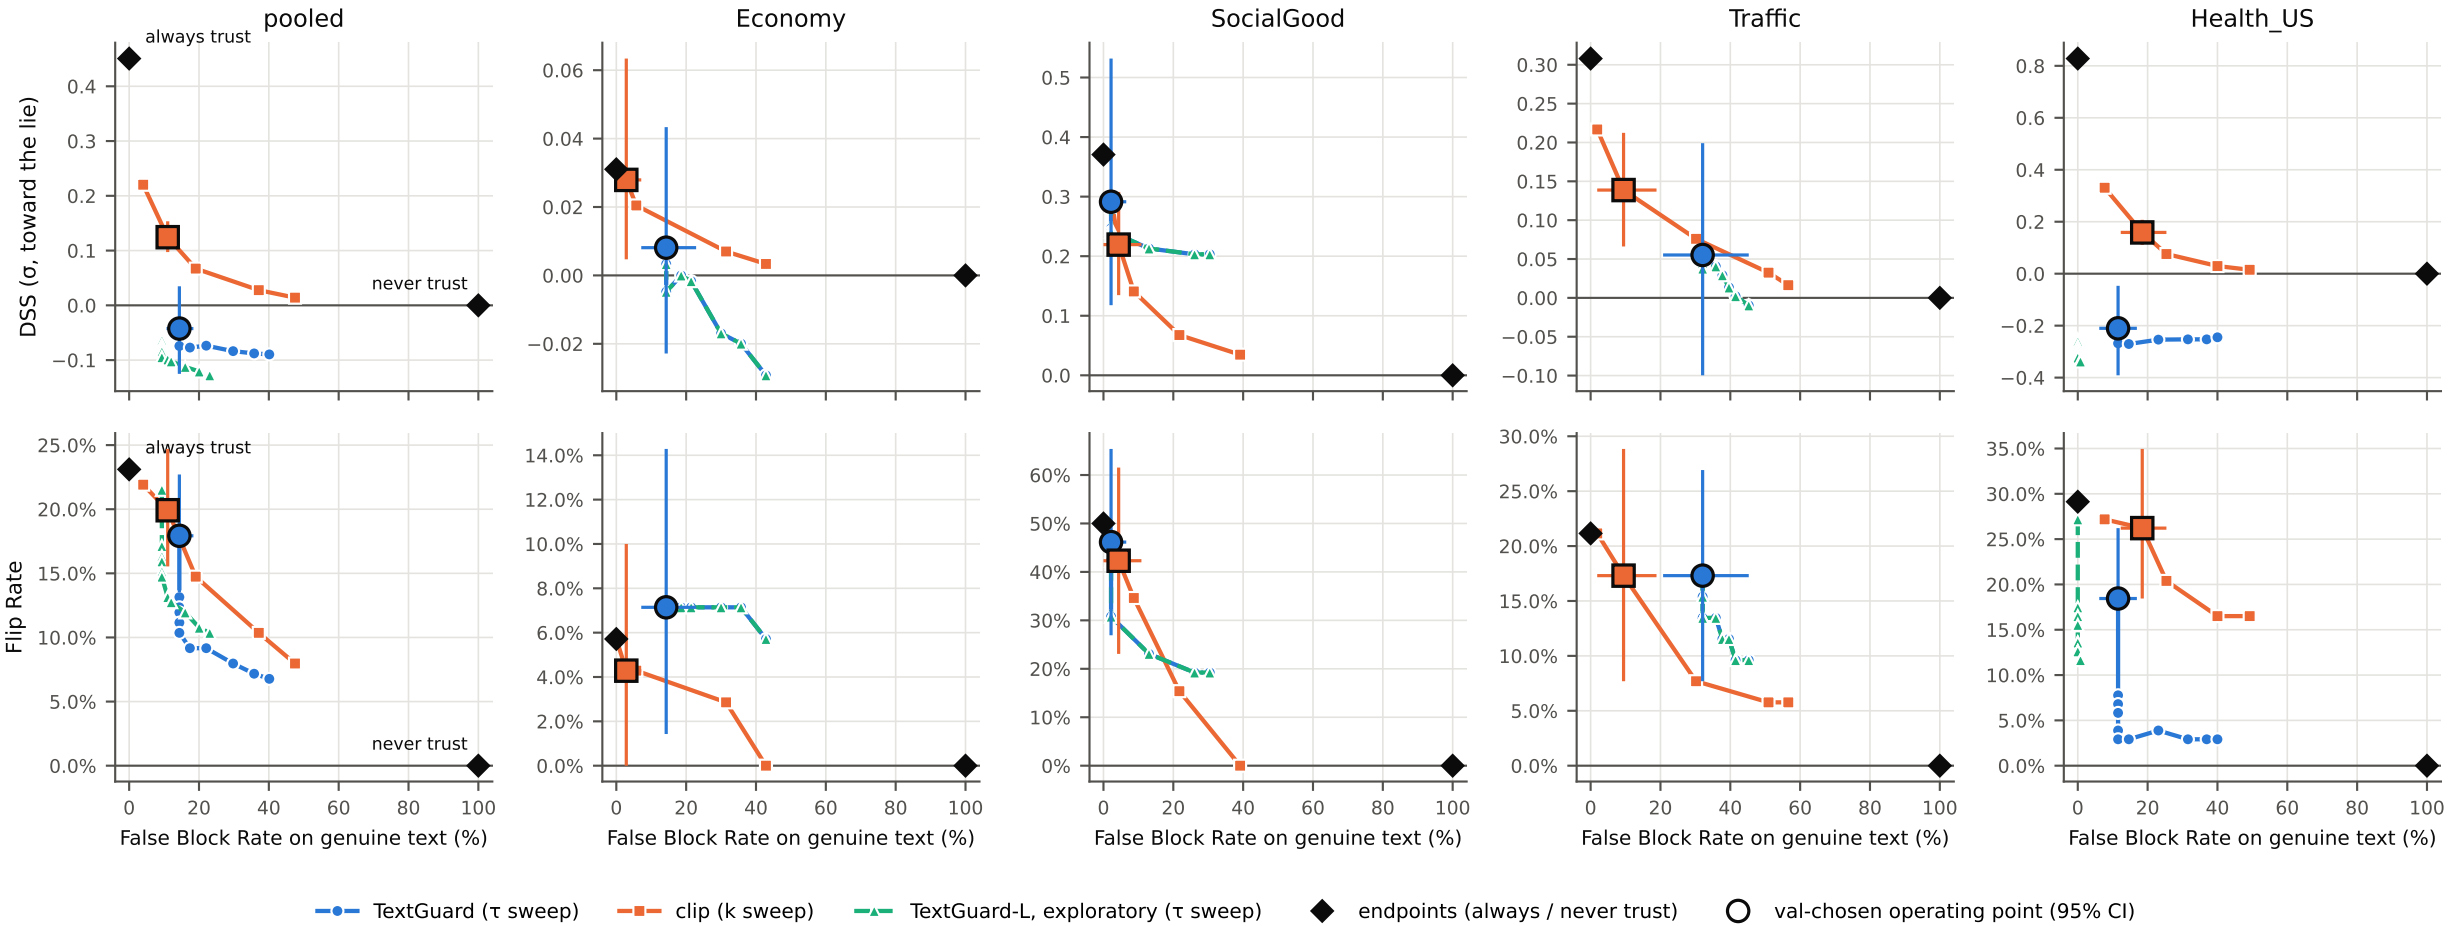
Model 5 (Qwen2.5-1.5B), flipped attack, test: manipulation vs genuine text blocked
pooled
Economy
SocialGood
Traffic
Health_US
always trust
never trust
DSS (σ, toward the lie)
Flip Rate
False Block Rate on genuine text (%)
TextGuard (τ sweep)
clip (k sweep)
TextGuard-L, exploratory (τ sweep)
endpoints (always / never trust)
val-chosen operating point (95% CI)